

Search.ai

Product Requirements Document (PRD) — v1.0

Tagline: Turn uncertainty into your next best move.

Document purpose: This PRD defines the complete product vision, user experience, functional requirements, AI/research system, comparison engine, safety controls, monetization, technical architecture, roadmap, and success measures for Search.ai. It is designed as a practical blueprint for a student founder building an MVP first, then expanding toward a production-grade decision-intelligence platform.

Product category: AI answer engine + research assistant + personalized tool/product/service comparison + action-planning workspace.

1. Product Summary

Search.ai is a decision-intelligence platform for people overwhelmed by too many choices. A user describes a problem in natural language—such as choosing study tools, launching a small business, finding software for a team, planning a career move, managing a budget, or building an app. Search.ai asks for the context that matters, researches relevant options, compares price and trade-offs, explains its recommendation, and produces an actionable plan.

The product should not position itself as “another chatbot.” Its core value is helping users make a defensible decision quickly: fewer tabs, transparent costs, credible alternatives, source visibility, and a next-step plan.

Problem statement

- People often open 10–30 tabs, read conflicting listicles, and still cannot decide which tool, service, path, or plan is right for their specific situation.
- Generic search results rank popularity or SEO performance, not personal fit, budget, skill level, location, risk tolerance, or time constraints.
- Most AI answers can be generic, unsupported, outdated, or overconfident; users need sources, trade-offs, and practical next steps.
- Students, freelancers, founders, small businesses, and everyday users frequently need low-cost, beginner-friendly options—not enterprise recommendations.

Product promise

Input: “I want to start an online business under \$10,000 and I have no coding experience.”

Output: A decision brief with a recommended path, top tools, price range, why the choice fits, alternatives, risks/trade-offs, sources, and a 7-day action plan.

2. Vision, Goals, Non-goals

Vision

Become the trusted place where a user can turn any important question into a clear, researched, context-aware decision.

MVP goals

- Help a user get a useful, structured recommendation for common problems across selected initial domains.
- Make recommendations transparent by showing fit reasons, cost, limitations, alternatives, and sources.
- Provide a simple action plan so research turns into execution.
- Create a reusable decision workspace where users can save, compare, share, and revisit decisions.
- Build a cost-controlled student-founder MVP before adding expensive real-time research features.

Non-goals for v1

- Do not claim professional medical, legal, tax, financial, or investment advice.
- Do not promise perfectly real-time pricing without a verification pipeline.
- Do not support every field with equal depth on day one; launch with high-value categories and expand deliberately.
- Do not build a social network, marketplace, full browser, or autonomous agent in the initial product.
- Do not use hidden sponsored ranking. Sponsored or affiliate relationships must be disclosed and cannot override relevance.

3. Users and Jobs-to-be-Done

Persona	Primary need	Example question	Value delivered
Student	Choose affordable learning and productivity tools	"How should I prepare for my exams?"	Free/low-cost stack + study plan
Freelancer	Find clients, organize work, invoice and build portfolio	"What tools do I need to start freelancing?"	Workflow + outreach + portfolio plan
Small business owner	Get customers and select affordable business tools	"How can my bakery get more orders?"	Local growth stack + execution steps
Founder	Choose an MVP stack and validate ideas	"How do I launch this app cheaply?"	Build-vs-buy comparison + cost/tech plan
Professional	Make work/process/software decisions	"Which CRM fits a 3-person sales team?"	Role-aware tool comparison
Job seeker	Choose learning, portfolio, application and career tools	"How do I move into frontend development?"	Skill-gap plan + job-search stack
Everyday user	Reduce confusion around common decisions	"How do I build a monthly budget?"	Simple, safe, practical roadmap

Key jobs to be done

- When I have too many possible tools or paths, help me identify the few options that fit my situation.
- When I am worried about cost, show what I will actually pay and what I give up at each budget level.
- When I receive a recommendation, explain why it was selected and what would make another option better.
- When I finish research, give me a concrete plan I can start today.
- When I return later, let me find my saved decisions, comparisons, notes, and follow-up tasks.

4. Core User Experience

Primary flow

Ask → Add context → Research → Compare → Decide → Act → Save/Share → Follow up.

Step	User action	System behavior	Output
1. Ask	Writes a problem naturally	Detects intent/domain and extracts decision criteria	Initial interpretation
2. Context	Chooses budget, experience, priority, location, timeframe	Asks only high-value follow-up questions	Personalized decision constraints
3. Research	Clicks "Map my options"	Retrieves trusted sources and tool data; normalizes evidence	Candidate options
4. Compare	Views ranked options	Scores fit and displays price, trade-offs and limitations	Comparison matrix

Step	User action	System behavior	Output
5. Decide	Selects Best Fit / Best Value / Alternative	Explains reasoning and uncertainty	Decision brief
6. Act	Uses generated next steps	Creates checklist, reminders or export	Action plan
7. Save	Saves result in workspace	Stores decision, sources and user notes	Reusable project record

Decision Brief: required output format

- **Your goal:** Restatement of the user's request and constraints.
- **Recommendation:** A concise best-fit path; never framed as universally best.
- **Confidence:** High / Medium / Low with an explanation of uncertainty.
- **Estimated cost:** Setup cost and recurring cost range where available.
- **Top options:** Usually 3 options: Best Fit, Best Value, Best Alternative.
- **Comparison:** Features, pricing, learning curve, setup time, region support, privacy considerations, pros, and cons.
- **Why this fits:** Explicit mapping from user constraints to recommendation.
- **Trade-offs:** What the user loses, risks, or must do manually.
- **Sources:** Direct source cards, source dates, and citation markers.
- **Action plan:** 3–7 actionable next steps, ideally time-boxed.
- **Follow-up chips:** Relevant questions that refine the decision.

5. Feature Inventory

Feature	MVP	Later	Description
Natural language problem input	Yes	—	Accept English, Hindi and Hinglish; provide examples and validation.
Decision context controls	Yes	—	Budget, skill level, location, deadline, priority, risk tolerance.
Domain detection	Yes	Improve	Classifies business, education, career, technology, finance, marketing, productivity, general.
Curated tool database	Yes	Expand	Structured records for selected categories; verified links and dates.
Rule-based fit scoring	Yes	Improve	Transparent weighted scoring based on stated criteria.
Best Fit / Value / Alternative	Yes	—	Three decision lenses, each with rationale and trade-offs.
Comparison matrix	Yes	Improve	Side-by-side product/tool comparison with filters.
Why-this-recommendation	Yes	—	Explainable recommendation reasons.
Source cards and citations	Yes	Improve	Official pricing/docs first; freshness date; evidence drawer.
Action plan generator	Yes	Improve	3–7 day or checklist-style plan.

Feature	MVP	Later	Description
Follow-up prompts	Yes	—	Contextual refinements without a new search.
Save decisions	Yes	—	Personal workspace, notes, bookmarks, statuses.
Shareable decision brief	Later	Yes	Read-only share link; privacy controls.
PDF export	Later	Yes	Export action plan/comparison with sources.
File/link analysis	Later	Yes	Analyze user-provided PDF, CSV, URL or document.
Live web research	Later	Yes	Search/retrieval with grounded citations and source selection.
Team workspace	Later	Yes	Shared research, comments, roles and audit trail.
Browser extension	Later	Yes	Compare tools/pages while browsing.
API	Later	Yes	Decision brief generation for third-party products.

6. Functional Requirements

6.1 Landing and acquisition

- Premium landing page with strong product narrative, interactive demo, clear pricing, FAQ, responsible-AI notice, and one primary CTA.
- Hero supports a problem input demo without requiring sign-up.
- A visual “Decision Core” can communicate the product story, but all content must remain readable without WebGL.
- Capture waitlist or sign-up intent with email and optional use-case selection.
- No fabricated customer logos, usage metrics, ratings, testimonials, or performance claims.

6.2 Authentication and onboarding

- Email magic link and Google sign-in for MVP.
- Onboarding collects preferred language, broad goals, country/region, skill level, default budget range, and selected use cases.
- User can skip onboarding and edit preferences later.
- Consent screen explains data handling, AI limitations, and optional use of personalization data.

6.3 Ask and clarify

- Question field supports 2,000+ characters, paste, and later voice input.
- System extracts: goal, category, budget, deadline, experience, location, constraints, preferred tools, and risk preferences.
- Ask no more than 1–3 follow-up questions per turn unless the user chooses “deep brief.”
- Allow the user to override auto-detected category and constraints.
- For sensitive/high-stakes categories, show boundary guidance and encourage professional support when appropriate.

6.4 Search and research

- Initial MVP uses curated tool records and controlled templates; label output as curated/demo where applicable.
- Production research retrieves from an allowlist of official vendor sites, documentation, pricing pages, public policies, and credible sources.

- Store source URL, title, publisher, retrieval time, excerpts, and claim linkage.
- Prefer primary sources for pricing, features, security claims, availability, and integrations.
- Show “last verified” or “retrieved on” metadata; do not imply all data is live if it is not.

6.5 Recommendation and ranking

- Return a compact top-three result by default; “show more options” available on demand.
- Rank using user constraints, hard eligibility filters, cost, required capability, ease of use, location support, and evidence confidence.
- Include an explanation for why a highly popular option was not selected.
- Separate organic ranking, sponsored results, and affiliate links. Sponsored results must have a visible label and cannot replace the top organic recommendation.
- Never invent price, integrations, user reviews, availability, certifications, or product features.

6.6 Comparison experience

- Compare up to 5 tools/services at a time.
- Columns: tool, best for, starting price, billing basis, free plan, setup time, learning curve, core features, limitations, region support, privacy/security notes, fit score, source date.
- User filters: free first, beginner-friendly, local support, no-code, privacy-focused, open-source, setup under one day.
- Allow user to pin comparison criteria and hide irrelevant rows.
- Explain missing data rather than filling blank values with assumptions.

6.7 Workspace and action

- Save a decision to a project/workspace.
- Add personal notes, completion status, tags, reminder date, and selected option.
- Generate checklists and time-boxed plans.
- Later: export PDF, send to Notion, Google Docs, calendar, Slack, or email—with explicit user confirmation before external writes.
- Provide decision history so users can revisit assumptions and compare past choices.

7. Initial Domain Scope

Domain	MVP questions	Core entities	Safety note
Business	Local growth, CRM, online store, customer tools	Marketing, messaging, storefront, analytics tools	No guaranteed revenue claims
Education	Study plans, notes, flashcards, courses	Learning tools, curricula, schedules	Not a substitute for teachers/exam rules
Career	Portfolio, skills, freelance workflow, job-search tools	Courses, portfolio tools, trackers	No employment guarantees
Technology	Build stack, hosting, database, no-code vs code	Frameworks, cloud, APIs, dev tools	Avoid unsafe code/security claims
Finance	Budgeting, tracking, personal organization	Budget apps, spreadsheets, reminders	General information only; no investment/tax advice
Marketing	Content, social scheduling, email, analytics	Design, social, email, measurement tools	No promised conversion outcomes

Domain	MVP questions	Core entities	Safety note
Productivity	Tasks, planning, focus, knowledge management	Task, calendar, note tools	Avoid health/diagnosis claims
General	Everyday research and decisions	Broad vetted categories	Ask clarifying questions first

Recommended launch order: Business, Education, Career, Technology, Productivity. These are high-frequency, lower-regulatory-risk domains and fit a student-founder's ability to curate quality tool data.

8. AI and Recommendation System

System pipeline

User query → intent/domain classifier → context extraction → clarifying questions → candidate retrieval → eligibility filters → fit scoring → evidence validation → structured decision brief → safety/policy review → UI rendering.

Fit-score model (explainable baseline)

Signal	Typical weight	Notes
Goal / use-case fit	25	Does tool solve the stated job-to-be-done?
Budget fit	20	Starting + recurring costs versus user budget.
Required capability fit	20	Must-have features and integrations.
Ease / skill-level fit	10	Learning curve, setup complexity, operational burden.
Location / availability fit	10	Country, language, payment and regional availability.
Time-to-value	5	Expected setup time and speed of first useful result.
Evidence confidence	5	Strength/freshness of supporting sources.
User preference fit	5	Privacy, open-source, team size, ecosystem preferences.

Scores should be presented as a decision aid, not scientific certainty. Display the criteria that influenced the score and let the user modify weights later.

Recommendation modes

Mode	Objective	Example
Best Fit	Highest overall weighted fit	Balanced low-cost beginner stack
Best Value	Maximum practical benefit per cost	Free-first stack with manual work
Best Alternative	Strong alternative with different trade-off	More capable paid option for future scale
Custom	User changes criteria weights	Privacy-first or fastest-setup mode

AI skills required in the product

- Intent classification and domain routing.
- Entity extraction: budget, country, time horizon, skills, product names, company names, goals, constraints.

- Question generation: ask only information that changes the recommendation.
- Retrieval-augmented generation (RAG): synthesize only from retrieved/curated evidence.
- Structured output generation: validated JSON for decision briefs.
- Comparison synthesis: convert normalized data into useful trade-offs.
- Plan generation: actionable, non-fabricated steps consistent with recommendation.
- Citation placement: attach source references to factual claims.
- Safety classification: identify medical/legal/financial/high-stakes topics and apply safe response policy.
- Language adaptation: explain in English, Hindi, Hinglish, and later regional languages without changing factual meaning.

9. Tool Knowledge Base

Tool record schema

Each recommended product/tool must exist as a structured record, not only as LLM text. This makes comparisons filterable, auditable, and maintainable.

Field	Description
id, slug, name, category	Stable identity and taxonomy
official_url, pricing_url, docs_url	Primary source links
description, use_cases	Human-reviewed short explanation
starting_price, currency, billing_cycle	Price fact plus source/retrieval date
free_plan, trial, price_notes	Eligibility and limitations
features, integrations, platforms	Normalized capabilities
setup_time_estimate, learning_curve	Practical adoption metadata
countries_supported, languages	Location fit
privacy/security notes	Only sourced, qualified statements
pros, tradeoffs, best_for, not_for	Decision-oriented editorial data
last_verified_at, confidence, sources	Freshness and evidence tracking
affiliate_disclosure, sponsored_status	Transparency metadata

Data operations

- Start with 50–100 curated tools across 5 launch domains.
- Use official product pricing and documentation pages as primary sources.
- Store verification dates; queue records for review when old or when price scraping detects a change.
- Add human review before publishing high-impact tool changes.
- Create a correction workflow: users can report an outdated price or fact.

10. Trust, Safety, and Responsible AI

Core principles

- **Transparency:** Explain why something is recommended and show relevant evidence.
- **Honesty about uncertainty:** Use low-confidence labels and request more context where evidence is weak.
- **User control:** User can change criteria, inspect sources, remove saved data, and opt out of personalization.
- **No deception:** Clearly label demo data, affiliate links, sponsored placements, and stale/unverified facts.
- **Privacy by default:** Collect minimal personal data and protect user decision history.
- **High-stakes boundaries:** Do not replace qualified professionals in medical, legal, tax, or investment contexts.

High-stakes policy

Category	Allowed	Not allowed	Required UI
Medical/health	General wellness education, reputable resources	Diagnosis, treatment plan, emergency instructions as substitute for care	Safety notice; encourage clinician/emergency service where relevant
Legal	General educational guidance and official resources	Legal conclusions, filing instructions presented as legal advice	Jurisdiction caveat; professional/legal aid referral
Finance/tax/investing	Budgeting education and general financial literacy	Personalized buy/sell recommendations, tax/legal advice, guaranteed returns	Risk disclaimer; verified official sources
Employment	Tool/career recommendations, resume organization	Guaranteed job outcomes or fabricated qualifications	Truthfulness reminder

AI output safeguards

- Use strict JSON schema validation; reject malformed model output.
- Ground factual tool details in structured database or retrieved evidence.
- Run claim checks: price, availability, compatibility, and policy claims must have an attached source or be marked uncertain.
- Prevent false credentials, fabricated citations, and invented user facts.
- Rate-limit expensive/deep research and protect against prompt injection in user-provided links/documents.
- Maintain audit logs for source retrieval, recommendation inputs, model version, and output version.

11. UX and Design Requirements

Design system

- Search.ai should feel like a decision-intelligence product, not a generic chat interface.
- Primary visual principle: chaos becomes clarity. Use an original “Decision Core” or data-convergence motif.
- Use cinematic dark mode, high-contrast typography, calm motion, and structured editorial layouts.
- 3D/WebGL is enhancement only; all core functions must work with CSS fallback and on mobile.
- Avoid excessive neon, generic robot imagery, fake dashboards, and card overload.
- Use accessibility-first color contrast, keyboard navigation, focus states, semantic landmarks, reduced-motion setting, and screen-reader labels.

Key screens

Screen	Purpose	Core elements
Landing page	Explain value and convert visitor	Hero, interactive demo, method, comparison proof, pricing, FAQ
Ask screen	Capture problem/context	Text input, constraints, clarifying questions, examples
Research progress	Explain wait/processing	Status steps, cancellable run, no fake “thinking” claims
Decision Brief	Main answer surface	Recommendation, confidence, costs, options, sources, action plan
Comparison view	Evaluate candidates	Sortable matrix, filters, pinned criteria, source dates
Workspace	Return and organize	Saved decisions, projects, notes, statuses, reminders
Tool directory	Browse curated data	Categories, filters, tool profiles, corrections
Settings	Control privacy and preferences	Language, default criteria, export/delete data, account

12. Pricing and Monetization

Pricing below is a proposed starting point, not a promise. Before public launch, calculate the cost of model tokens, web research, storage, support, payment fees, and taxes. Keep plan limits explicit; expensive research must be metered.

Plan	Proposed price	Target user	Included capabilities
Free	■0/month	Students and casual decision-makers	10 standard questions/day; 3 deep comparisons/month; basic citations; 10 saved decisions
Plus	■299/month or ■2,499/year	Freelancers and regular users	Unlimited standard questions; 100 deep comparisons/month; detailed comparison; file/link analysis; unlimited saves; PDF export
Pro	■799/month or ■6,999/year	Founders and professionals	500 deep comparisons/month; advanced reports; private projects; priority processing; budget/market-fit scenarios
Teams	■1,499/seat/month	Small teams	Shared workspaces, comments, team shortlists, admin controls, usage analytics
Enterprise	Custom	Larger organizations	SSO, audit logs, custom data sources, SLA, procurement/security review

Revenue streams

- Subscriptions: core recurring revenue.
- Affiliate commissions: only with conspicuous disclosure and no ranking override.
- Teams/enterprise: collaborative research, security, integrations.
- API: structured decision brief/recommendation service for third parties.
- Sponsored placements: only as an explicitly separate, labeled section; never disguised as an organic result.
- Student plan: verified discount to increase adoption and reduce churn.

13. Technical Architecture

Student-friendly MVP stack

Layer	MVP choice	Why	Later evolution
Frontend	Next.js + TypeScript + Tailwind	Single codebase, fast UI development, deployable on Vercel	Native mobile/PWA if needed
UI	shadcn/ui + Framer Motion	Accessible primitives; controlled visual polish	Design system package
Auth	Supabase Auth	Fast email/Google auth	Enterprise SSO
Database	Supabase Postgres	Structured relational data, RLS, SQL, low setup	Managed Postgres with replicas/partitioning
Storage	Supabase Storage	Upload user files later	Dedicated object storage
Backend	Next.js route handlers/server actions	Simple MVP boundary	Dedicated services for retrieval/jobs
AI	Structured LLM output + curated data	Controlled cost and explainability	Model routing, caching, evals
Research	Curated DB first; controlled web search later	Avoid weak/unverified facts	RAG pipeline + source trust scoring
Jobs	Cron/scheduled functions	Refresh tool data, reminders	Queue/workers for scale
Observability	Structured logs + error tracking	Debug outputs and cost	Tracing, analytics warehouse

Logical architecture

Client UI → API gateway → decision orchestrator → (context extractor, candidate retriever, ranking engine, source verifier, LLM structured-response generator, policy guardrail) → Postgres/tool knowledge base + source store + user workspace. Background jobs refresh tool data, expire cache entries, generate exports, and schedule reminders.

Important engineering decisions

- Keep structured tool facts outside the LLM prompt whenever possible; use the model for interpretation and synthesis, not as the source of truth.
- Use JSON schema / Zod validation for all model responses before showing them to users.
- Cache repeated tool facts and public research to control cost; cache must preserve verification date.
- Use row-level security so each user only accesses their own decisions, files, and workspace.
- Separate analytics events from personally identifiable content where possible.
- Use feature flags for expensive research, file analysis, and experimental ranking logic.

14. Data Model

Table	Key fields	Purpose
users / profiles	id, name, locale, country, onboarding state	Account and preference profile
projects	id, user_id, title, status, tags	Workspace grouping
decisions	id, user_id, project_id, query, constraints JSON, status, selected option	Core decision record

Table	Key fields	Purpose
decision_options	id, decision_id, tool_id, rank, fit_score, rationale, tradeoffs	Ranked options
tools	id, slug, name, category, editorial metadata	Canonical tool record
tool_prices	tool_id, region, currency, amount, cycle, source_id, verified_at	Historical/region price facts
sources	id, url, publisher, title, retrieved_at, trust tier	Evidence records
claims	decision_id, option_id, claim text, source_id, confidence	Citation mapping
action_plans	decision_id, steps JSON, status, due dates	Execution plan
saved_items	user_id, tool_id/decision_id, note, tags	Bookmarks
feedback	user_id, decision_id, rating, correction, selected outcome	Quality learning loop
usage_events	user_id, event, plan, cost estimate	Quota, product analytics, billing
audit_logs	actor, action, resource, metadata, timestamp	Debug/security accountability

15. APIs and Structured Contracts

Core endpoints

Endpoint	Method	Purpose
/api/decisions	POST	Create decision request and return/retrieve decision brief
/api/decisions/:id	GET	Fetch decision and sources
/api/decisions/:id/refine	POST	Apply follow-up constraints/question
/api/decisions/:id/select	POST	Save chosen option
/api/projects	GET/POST	List/create workspaces
/api/tools	GET	Browse/filter tool directory
/api/tools/:slug	GET	Tool detail with freshness metadata
/api/compare	POST	Compare selected tools and criteria
/api/exports/:id	POST	Generate export after explicit request
/api/feedback	POST	Record quality/correction feedback

Decision brief response contract

Return validated structured data similar to: decision id, normalized query, constraints, confidence, recommendation, options[], fit reasons[], tradeoffs[], estimated_cost, sources[], action_plan[], follow_up_questions[], safety_notices[], data_freshness. The frontend renders this schema; it should not parse arbitrary prose.

16. Quality, Evaluation, and Analytics

North-star metric

Decision completion rate: percentage of started decision sessions in which a user saves a decision, selects an option, starts an action plan, or reports meaningful progress.

Product metrics

Area	Metrics
Acquisition	Landing-to-question conversion, sign-up conversion, waitlist conversion, source of traffic
Activation	First decision completed, first saved decision, first comparison opened, first action step checked
Quality	User rating, correction rate, citation coverage, stale-data rate, "not useful" reasons
Recommendation	Option click-through, selection rate, alternative-tab usage, follow-up rate
Retention	Weekly active decision makers, saved workspace return rate, project completion rate
Business	Free-to-paid conversion, ARPU, research cost per active user, gross margin
Safety	High-stakes routing rate, unsupported claim rate, policy escalation rate, user reports

Evaluation set

- Create 50–100 realistic benchmark queries across launch domains.
- For each query define expected constraints, acceptable options, unsafe answers, required caveats, and expected sources.
- Evaluate: factual grounding, price accuracy, relevance, explanation quality, actionability, language quality, and safety.
- Run regression tests when prompt, model, ranking weights, or tool data changes.
- Use human review for a sample of results before claiming quality publicly.

17. Security and Privacy

- Use HTTPS, secure cookies, protected server-side secrets, and environment variables; never expose provider API keys in client code.
- Enable Supabase Row-Level Security on every user-owned table.
- Provide account deletion, data export, and saved-decision deletion controls.
- Minimize uploaded files; encrypt at rest through provider controls and use signed URLs.
- Scan/limit file types and sizes before file analysis; do not execute user-provided code.
- Protect against prompt injection from webpages/files: treat fetched text as untrusted data, isolate instructions, and restrict tool permissions.
- Rate-limit abuse, use CAPTCHA only if needed, and throttle expensive deep-research endpoints.
- Do not store payment cards directly; use a payment processor such as Stripe/Razorpay in later stages.
- Maintain a privacy policy and clearly state retention/deletion behavior before public launch.

18. Phased Roadmap

Phase	Duration	Deliverable	Definition of done
0. Foundation	2–4 days	Brand, landing page, demo data, design system	Interactive demo works; no false claims; mobile usable

Phase	Duration	Deliverable	Definition of done
1. Decision MVP	1–2 weeks	Ask flow, 5 domains, curated tools, rule-based scoring, decision brief	User gets structured results and saves decisions
2. Workspace	1 week	Auth, projects, history, notes, action plans	Returning user can manage decisions
3. Research beta	2–3 weeks	Controlled web retrieval, citations, freshness, source cards	Outputs grounded with source metadata
4. Personalization	1–2 weeks	Preference weights, refined comparisons, feedback loop	Results change transparently with context
5. Monetization	1 week	Plans, quotas, billing integration, student discount	Usage metered and checkout tested
6. Team/Pro	Later	Sharing, exports, file/link analysis, team workspace	Permissions and privacy controls tested

First 7-day student-founder plan

- **Day 1:** Create new Git repository; define 5 launch categories and 50 tool records in a spreadsheet/JSON.
- **Day 2:** Build landing page and interactive non-AI Decision Brief demo.
- **Day 3:** Build Ask page, context selectors, and deterministic category detection.
- **Day 4:** Implement tool filtering, fit scoring, Best Fit/Value/Alternative tabs.
- **Day 5:** Build comparison table, sources/freshness UI, and action-plan templates.
- **Day 6:** Add Supabase auth and save-decision workspace.
- **Day 7:** Test with 10 users, collect feedback, fix top usability issues, and record a short demo video.

19. Risks and Mitigations

Risk	Why it matters	Mitigation
Outdated pricing	Breaks trust and decisions	Show verification date/source; refresh jobs; allow reports; avoid “live” claim without live data
Hallucinated facts	AI may invent features or prices	Structured DB, retrieval grounding, source-required claims, schema validation
High AI cost	Student budget can be exhausted	Curated MVP, quotas, caching, concise prompts, gated deep research
Generic recommendations	User sees no unique value	Collect context; explain fit; show trade-offs; personalize filters
Scope explosion	“Any field” becomes unbuildable	Launch with 5 domains; define domain playbooks; expand from evidence
Affiliate conflict	Users distrust rankings	Disclosure, separate sponsored module, documented ranking policy
Poor mobile performance	3D scenes can slow devices	Progressive enhancement, fallback, lazy loading, reduced motion
Sensitive advice	Health/legal/finance risk	Policy routing, disclaimers, primary sources, no professional claims
Privacy failure	Decision history can be sensitive	RLS, data minimization, deletion/export controls, secure storage

20. Acceptance Criteria for MVP

- A new visitor can submit a natural-language problem and receive a structured demo decision brief in under 5 seconds in mock mode.
- User can choose/edit budget, experience, location, and priority before generating a result.
- System supports at least 5 domains and 50 curated tool records.
- Every recommendation shows why it fits and at least one trade-off.
- Every displayed factual tool price/feature has a source URL and verification date, or is explicitly labeled illustrative/demo data.
- User can compare at least 3 tools side by side and filter by free/beginner-friendly criteria.
- User can save a decision after authentication and see it in a workspace.
- Sensitive-domain query shows the correct limitation/safety notice.
- Mobile layout works on a 360px-wide screen and core functions work without WebGL.
- No API secret, user data, or private file is exposed in client source code or Git history.
- Landing page clearly communicates who the product is for, what it does, pricing approach, and primary CTA.

21. Future Vision

- A personal “Decision OS” that remembers a user’s goals, constraints, tools, projects, and past choices with explicit permission.
- Collaborative team decision rooms for evaluating vendors, software stacks, hiring tools, project approaches, and budgets.
- Verified price-change alerts and “better alternative available” notifications.
- Deep vertical playbooks for Indian students, local businesses, freelancers, startup founders, and global remote professionals.
- Browser extension that converts any pricing page or product comparison into a user-specific decision brief.
- Explainable recommendation API for marketplaces, educational platforms, and business-software directories.

22. Final Product Principle

Search.ai should maximize clarity per minute. It should not make users read more, browse more, or trust an unexplained black box. The product wins when a user can say: “I understand my options, I know the trade-offs, and I know what to do next.”

For the first release, build the important part well: accurate curated data, explicit criteria, transparent comparisons, useful action plans, safe boundaries, and a clean interface. Add real-time AI research only after the deterministic product workflow is reliable.